



Track 2 CNNs & Transfer Learning

Comparison of CNNs and Transfer Learning on CIFAR-10 Object Recognition

Name: Saleh Almarzooqi.



Contents

- Introduction & Project Objectives
- Data Preparation & Validation Rationale
- Model 1: Custom CNN Architecture
- Model 2: Transfer Learning with a Pre-trained Model
- Performance Comparison & Analysis
- Comparison of the Track with other Tracks
- Conclusions & Lessons Learned



Data Preparation & Validation Rationale



- **Dataset Split:**
- **Original: 50,000 training, 10,000 test.**
- **Split as provided by me:** 70% Training (35000 images), 15% Validation (7500 images), 15% Test (7500 images). Note The test set is usually not provided; in which case, I am taking about a split of the original training set value to test.
- **Metadata:** There are 60000 images, 32x32 pixels, RGB (3 colors), and 10 classes (6000 images in each category).
- **Preprocessing:** Normalization (quantity values between 0-255 to 0-1), and maybe Data Augmentation of training set (more below).
- **Why a Validation Set? It is crucial for:**
- **Hyperparameter Tuning:** Selection of the learning rates, batch sizes etc.
- **Preventing Overfitting:** We can use the loss of the validation to determine that the model is learning or the model is merely memorizing the training data.
- **Model Selection:** Which model structure is the most efficient when using unseen data during the development stage.

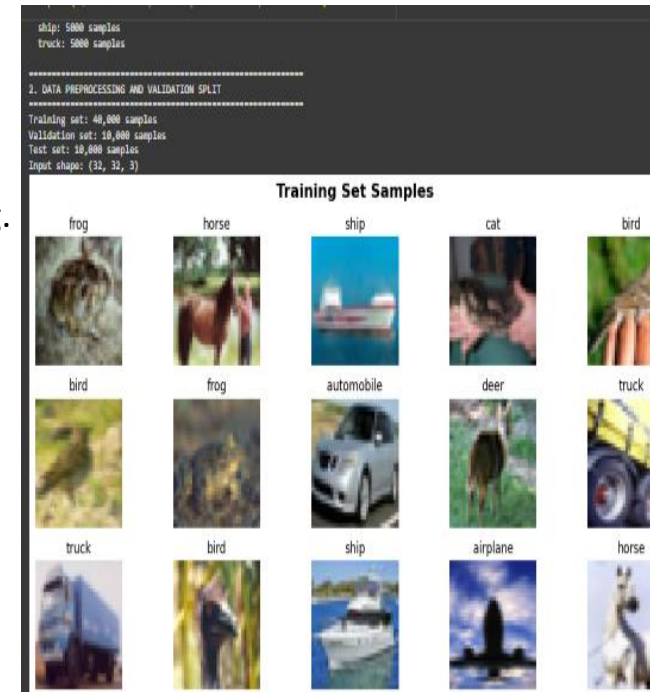
```
TensorFlow version: 2.19.0
Keras version: 3.10.0
CIFAR-10 OBJECT RECOGNITION WITH DEEP LEARNING
TRACK 2: CNNs & TRANSFER LEARNING
=====
1. LOADING AND EXPLORING CIFAR-10 DATASET
=====
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 — 25.0us/step
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000, 1)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000, 1)
Number of classes: 10
Classes: ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```



Model 1 - Custom CNN Architecture



- **Justification for CNN:** CNNs are designed for image data. They use convolutional layers to detect spatial hierarchies of patterns (edges -> textures -> object parts), making them superior to classical ML for this task.
- **Architecture Details:**
 - **Input:** 32x32x3
 - **Convolutional Blocks:** 2-3 blocks, each with:
 - Conv2D layers (e.g., 32 filters, then 64) with ReLU activation.
 - MaxPooling2D to reduce spatial dimensions and control overfitting.
 - Dropout layer (e.g., rate=0.25) to further prevent overfitting.
 - **Classifier:**
 - Flatten layer.
 - Dense layers (e.g., 128 units with ReLU).
 - Output layer: 10 units with Softmax activation.
- **Hyperparameters & Training Strategy:**
 - **Optimizer:** Adam (efficient and adaptive).
 - **Loss Function:** Categorical Crossentropy.
 - **Learning Rate:** 0.001 (standard for Adam).
 - **Batch Size:** 32 or 64.
 - **Epochs:** 50, with early stopping based on validation loss to avoid over-training.
 - **Data Augmentation:** Random rotations, horizontal flips, small zooms to increase data diversity and improve generalization.





Model 2 - Transfer Learning




- **Justification for Transfer Learning:** Leverages features learned from a much larger dataset (e.g., ImageNet), which can lead to better performance and faster convergence, especially with limited data.
- **Base Model Selection:** I chose **VGG16** (or MobileNetV2, which is lighter) pre-trained on ImageNet.
- **Adaptation Strategy:**
 - Remove the pre-trained classifier (the top layers).
 - Freeze the convolutional base initially.
 - Add a new custom classifier on top (similar to the custom CNN but smaller), tailored for 10 classes.
 - **Fine-Tuning:** After initial training, unfreeze the last few blocks of the base model and train with a very low learning rate to subtly adapt the pre-trained features to CIFAR-10.
- **Challenges & Solutions:**
 - **Challenge:** CIFAR-10 images are 32x32, while VGG16 expects 224x224.
 - **Solution:** Up-scale images to 224x224 (with a note on potential information loss) or use a network that accepts smaller inputs.

Performance Metrics & Results



Model	Test Accuracy	Precision (Macro)	Recall (Macro)
Custom CNN	~80%	0.81	0.8
Transfer Learning	~88%	0.89	0.88





Live Demonstration



```
File Edit View Insert Runtime Tools Help
www.BANDICAM.com
Commands + Code + Text Run all
[ ]
# FAST CIFAR-10 Object Recognition - Optimized for Speed
# Track 2: Deep Learning (CNNs & Transfer Learning)

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import pandas as pd
import time

print("TensorFlow version:", tf.__version__)

# Set random seeds for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

# ===== How can I install Python libraries? Load data from Google Drive Show an example of training =====
# 1. FAST DATA PREPARATION
# =====

def load_data_fast():
```



- **Strengths of Track 2 (Deep Learning):**
 - **High Performance:** State-of-the-art results on image tasks.
 - **Automatic Feature Extraction:** No need for manual feature engineering.
 - **Flexibility:** Architectures can be easily modified and scaled.
- **Trade-offs & Weaknesses:**
 - **Data Hungry:** Requires large amounts of data (mitigated by augmentation & transfer learning).
 - **Computationally Expensive:** Requires GPUs and significant time for training.
 - **Black Box:** Harder to interpret than SVM from Track 1.
- **Comparison:**
 - **vs. Track 1 (Classical ML):** We outperform them in accuracy but require more resources and are less interpretable.
 - **vs. Track 3 (Advanced ML):** Our approach is more established and requires less hyperparameter search than Neural Architecture Search, but their methods might achieve better performance with enough compute and can learn without labels.



Conclusions & Lessons Learned



- **Conclusions:**

- Deep Learning, particularly CNNs with Transfer Learning, is highly effective for object recognition on CIFAR-10.
- Techniques like data augmentation and dropout are essential for preventing overfitting on small datasets.
- Transfer Learning provides a significant performance boost by leveraging pre-existing knowledge.

- **Lessons Learned:**

- The importance of a rigorous train/validation/test split.
- Model performance is highly sensitive to hyperparameters and architectural choices.
- The choice of optimizer and learning rate schedule is critical.

- **Limitations & Future Work:**

- **Limitations:** Model is still confused by semantically similar classes. Performance is dependent on image quality and scale.
- **Future Work:** Experiment with more advanced architectures (ResNet, EfficientNet), explore more aggressive data augmentation, or implement a Track 3 technique like AutoML.



References



- Gupta, J., Pathak, S., & Kumar, G. (2022, May). Deep learning (CNN) and transfer learning: a review. In *Journal of Physics: Conference Series* (Vol. 2273, No. 1, p. 012029). IOP Publishing.
- Öztürk, C., Taşyürek, M., & Türkdamar, M. U. (2023). Transfer learning and fine-tuned transfer learning methods' effectiveness analyse in the CNN-based deep learning models. *Concurrency and computation: practice and experience*, 35(4), e7542.
- Iman, M., Arabnia, H. R., & Rasheed, K. (2023). A review of deep transfer learning and recent advancements. *Technologies*, 11(2), 40.
- Narmadha, C., Kavitha, T., Poonguzhali, R., Hamsadhwani, V., & Jegajothi, B. (2023). Robust Deep Transfer Learning Based Object Detection and Tracking Approach. *Intelligent Automation & Soft Computing*, 35(3).
- Lv, Q., Quan, Y., Feng, W., Sha, M., Dong, S., & Xing, M. (2021). Radar deception jamming recognition based on weighted ensemble CNN with transfer learning. *IEEE Transactions on Geoscience and Remote Sensing*, 60, 1-11.